

In the last three lectures we have seen:

- Elixir concurrency (send, receive, ...)
- Elixir fault handling (links, monitors, ...)
- Elixir distribution (node-node communication)

We have seen examples of Elixir programs using these primitives to implement rather complex behaviour in an elegant and compact fashion

- So, everything is perfect!
- Or? ...

Concurrency and Distribution

Let us compare with non-concurrent programming in Elixir:


- **Everything** can be programmed using just recursion
- **But** to prevent re-inventing code (and re-inventing bugs!) we prefer to use “higher-level” libraries like *Enum*, *List*, etc

What are the “higher-level” libraries for concurrency in Elixir?


- Today: from Erlang – OTP: Open Telecom Platform:
 - ▶ *GenServer* – a library for client/server applications
 - ▶ *Supervisor* – a library to program fault-tolerant applications
- Later lectures:
 - ▶ *Agent*, *Task*, *PubSub*, *GenStage*, ...

The *GenServer* – Generic Server – Library

- Generic Servers provide a standard way to implement client–server interactions
- One server interacts with an arbitrary number of clients:




The *GenServer* – Generic Server – Library



- Clients makes *requests* to the server, which optionally *replies*
- Clients *block* until the server replies, or make *non-blocking* calls
- A server has a state, which is preserved between requests
- A server handles one call at a time

The *GenServer* – Generic Server – Library



- Clients makes *requests* to the server, which optionally *replies*
- Clients *block* until the server replies, or make *non-blocking* calls
- A server has a state, which is preserved between requests
- A server handles one call at a time
 - ▶ however a server may postpone the “return” from a call until later, thus permitting to handle later calls first
 - ▶ Example: a server implementing a “counter” with initial value 0, and with two operations:
 - ★ `inc` – never blocks
 - ★ `dec` – blocks until the value of the counter is positive

Accessing a GenServer from a Client

- Creating a Generic Server defined by the module *MyServer*:

```
# Start a GenServer registered as MyServer  
GenServer.start_link(MyServer, InitArg, name: MyServer)
```

- A blocking call to the server:

```
# request can be any value  
# timeout in milliseconds for when call fails  
result = GenServer.call(MyServer, request, timeout)
```

- A non-blocking call to the server (no return value):

```
GenServer.cast(MyServer, request)
```

Defining a GenServer

The actions of the GenServer is defined in a number of functions that the user must define:

```
defmodule MyServer do
  use GenServer    # Declare that the module uses
                  # the GenServer library

  def init(InitArg)    # Returns the initial state
                      # of the server

  def handle_call(...) # Handles a blocking call
  def handle_cast(...) # Handles a non-blocking call
end
```

Handling Blocking Calls

```
def handle_call(request, from, state) do
  ...
end
```

- request is the argument from *GenServer.call*
- from represents the caller
- state is the state when the call was made

Handling Blocking Calls: Values Returned

```
def handle_call(request, from, state) do
  ...
  {:reply, return_value, new_state}
end
```

- The `handle_call` function *always* returns a new state, and *may* return a value to the caller, thus unblocking it
- Concretely a tuple is returned, either:
 - ▶ `{:reply, return_value, new_state}` – returns a new state and replies with the value `return_value` to the caller
 - ▶ `{:noreply, new_state}` – just returns a new state
- Note that when handling a new call, it is possible to reply to a previous call by calling `GenServer:reply(previous_from, return_value)`

Handling Non-Blocking Calls

- Non-blocking calls are handled in `handle_cast`:

```
def handle_cast(request, state) do
  ...
  {:noreply, new_state}
end
```

- The caller cannot receive a return value

Example: A Resource Handler

- A process *pidH* controls some piece of sensitive hardware (e.g., the brakes in a car)
- To interact with the hardware (e.g., brake the car) other processes sends messages to *pidH*
- To prevent bad interactions between different processes interacting with the hardware only one process may communicate with *pidH* at any given time.
- A client (process) must use a simple access protocol before sending messages to *pidH*:

```
Hw.request()           # request access to the hardware,  
                        # return when permission given
```

```
send(pidH, msg1)       # interact with hardware  
...  
send(pidH, msgN)
```

```
Hw.release()          # notify stopping access to hardware
```

- Our task is to use a GenServer to implement the protocol, or in other words, to implement the *Hw* module.

The Server State

- Multiple clients may request access to the hardware at the same time, but we should only grant access to a single process at a time
- As request and release messages will arrive at the server in any order, it has to manage a queue of hardware access requests
- What should the server state be?
- One option: let the state be a list [*Head* | *Rest*] where
 - ▶ *Head* is the client (process) currently accessing the resource
 - ▶ and *Rest* are the clients (processes) waiting to access the resource

Handling Client Errors

- What happens if a client fails crashes while accessing the resource? (after request but before release)
- The GenServer will stay “locked” forever, other clients in the queue will never get access
- How to fix?
- We could monitor/link to the client and removing it from queue when it crashes
 - ▶ We have to implement the following callback function:

```
handle_info(msg, state) do
  ...
  {:noreply, newstate}
end
```

which is called when non-call and non-cast messages are received by the *GenServer* process


Error Handling in the Generic Server component

Note some nice properties of error handling in generic servers:

- We only handle errors in one place in the code
- We only handle errors at very controlled points in time (when not processing a request)
- We control error handling – we do not letting error handling control us!


The Supervisor Component

- Elixir Applications are often structured as *supervision trees*, consisting of *supervisor* and *worker* processes




- A supervisor
 - ▶ starts child processes in a specified order
 - ▶ monitors them – and possibly restarts them if they fail
 - ▶ stops them on request
- The actions of the supervisor are described in a declarative fashion (as a text description)
- Note that a child process may itself be a supervisor

Supervision Dynamics



- We assume that the children has been started in the order C1, C2, C3.
- When a child process C2 dies, its supervisor S3 is notified.
If the child C2 is declared
 - ▶ **:permanent** – the child is always restarted
 - ▶ **:temporary** – the child is not restarted (even if it terminated abnormally)
 - ▶ **:transient** – the child is restarted only if it terminated abnormally
- If the decision is that C2 should be restarted then different sibling restart strategies can be specified in supervisor S3.
 - ▶ **:one_for_one** – only restart the failing child, not its siblings
 - ★ i.e., only restart C2, do not restart C1 and C3
 - ▶ **:one_for_all** – restart **all** siblings
 - ★ i.e., terminate first all its siblings C1 and C3, and then restart C1, C2 and C3
 - ▶ **:rest_for_one** – terminate all siblings started **after** the failed child
 - ★ i.e., since C3 was started after C2 terminate it, and then restart C2 and C3



- One can also control the frequency of restarts, and the maximum number of restarts to attempt – it is no good having a process continuing to restart and crash:
 - ▶ `:max_restarts` – the maximum number of restarts permitted in a time frame. If exceeded the supervisor terminates *abnormally* (default 3).
 - ▶ `:max_seconds` – the time frame (default 5 seconds).

Supervisors: under-the-hood

- A child process being supervised by a supervisor should be *linked* to the supervisor
- Read more about supervisors on the web:
 - ▶ <https://hexdocs.pm/elixir/1.13/Supervisor.html>
 - ▶ and lots of tutorials...

Example: Worker Jobs

To experiment with supervisors we program a simple executor of jobs

```
Jobs.execute(  
  %{  
    job1_name: f_1,  
    ...  
    jobN_name: f_n  
  })
```

- A job has a name and a execution specification, i.e., a function with arity zero.
- The *Jobs*.execute function starts each job in a separate process, in parallel, collects the results, and returns a map where the keys are the job names and the values the job results.
- Example:

```
Jobs.execute(  
  %{ time: fn () -> Time.to_string(Time.utc_now()) end,  
    date: fn () -> Date.to_string(Date.utc_today()) end})
```

==>

```
%{ time: "08:41:41.206419", date: "2022-04-22" }
```

Idea for implementing `execute(jobs)`:

```
def execute(jobs) do
  # Spawn collector process
  myPid = self()
  spawn(fn () -> register(self(), :collector)
        collector(myPid, map_size(jobs)) end)


  # Spawn job processes
  Enum.each(jobs, fn {job, f} ->
    spawn(fn () -> send(:collector, {job, f.()}) end)
  end)

  # Wait for results from collector process
  receive do x -> x end
end
```


- 1 When `Jobs.execute(...)` is invoked from the “shell” a “collector process” is started which is responsible for collecting the results from the jobs (time, date)
- 2 The job processes (time, date) start, compute their results, and sends them to the collector
- 3 When the collector process has received both results, it sends them to the “shell”

Supervision of the Job executor

- How do we program a *robust* Job executor?
 - ▶ One process gets started for each job to execute – *J1, J2, ...*
 - ▶ Additionally we create one process to accumulate the results – *Collector*
- Concretely we use a supervision tree of two three levels: one supervisor *JobsSupervisor* for the all the job workers, and one supervisor *AppSupervisor* for the *Collector* process **and** for the worker supervisor:




Decision: restart strategy for a Job worker process




- What should we do if a worker process terminates *abnormally*?
 - ▶ restart it
- What should we do if a worker process terminates *normally*?
 - ▶ nothing
- So a worker is a **:transient** process

Decision: restart strategy for the *Collector* process




- Similarly the *Collector* process may terminate normally, it should be restarted only when it terminates abnormally
- It is a **transient** process

Decision: restart strategy for the *AppSupervisor*




- What should we do if the *JobsSupervisor* process terminates abnormally?
 - ▶ restart it **and** also restart the *Collector* process
- What should we do if the *Collector* process terminates abnormally?

Decision: restart strategy for the *AppSupervisor*



- What should we do if the *JobsSupervisor* process terminates abnormally?
 - ▶ restart it **and** also restart the *Collector* process
- What should we do if the *Collector* process terminates abnormally?
 - ▶ restart it, **and** stop and restart the *JobsSupervisor* (and all the jobs – happens automatically)

Decision: restart strategy for the *AppSupervisor*



- What should we do if the *JobsSupervisor* process terminates abnormally?
 - ▶ restart it **and** also restart the *Collector* process
- What should we do if the *Collector* process terminates abnormally?
 - ▶ restart it, **and** stop and restart the *JobsSupervisor* (and all the jobs – happens automatically)
- So the *AppSupervisor* should use the strategy `one_for_all`

Starting Supervisors: example

To start the top supervisor process, executing jobs, we invoke:

```
Supervisor.start_link(children,  
                        strategy: :rest_for_all,  
                        name: AppSupervisor)
```

- strategy specifies the restart strategy for the children
- name specifies a registered name for the supervisor process
- children is a list of modules and arguments for the supervised children (collector, jobsSupervisor):

```
children =  
  {Collector, [self(), map_size(jobs)]},  
  {JobsSupervisor, [jobs]}
```

- ▶ For each child, the supervisor
 - 1 retrieves its child specification,
 - 2 and then starts the child (see next slide)
- ▶ self() is passed as a parameter to the collector so that it knows where to send the result
- ▶ map_size(jobs) is passed as a parameter so that it knows how many results to expect

Collector code

- To create the Collector from the recipe `{Collector, arg}`, the supervisor first retrieves a “child” specification by calling `Collector.child_spec(arg)`:

```
def child_spec([parent_pid, num_jobs]) do
  %{
    id: Collector,          # identity
    restart: :transient # restart only if abnormal exit,
    start: {Collector, :start_link, [parent_pid, num_jobs]}
  }
end
```

- The start key specifies how to start the child (Collector), a tuple (module, function, arguments). The function should start and link to a new process, and return a tuple `{:ok, pid}` where `pid` is the pid of the new process:

```
def start_link(parent_pid, num_jobs) do
  {:ok,
   spawn_link(fn () ->
     Process.register(self(), :collector)
     # call main service loop
     collector(parent_pid, %{}, num_jobs)
   end)}
end
```

Supervisors: advantages

- A nice way for specifying the termination “dynamics” of an application
 - ▶ What should the application do when one of its components fails?
 - ▶ And what should it do when one of its components terminates normally?
- Permits starting and stopping the entire application with one just command
 - ▶ All the processes in the application are started when the topmost supervisor is started
 - ▶ All the processes in the application are stopped when the topmost supervisor is stopped